

# Lab 3 Report:

## MNIST Classification with FCN

Name: Forest Tschirhart

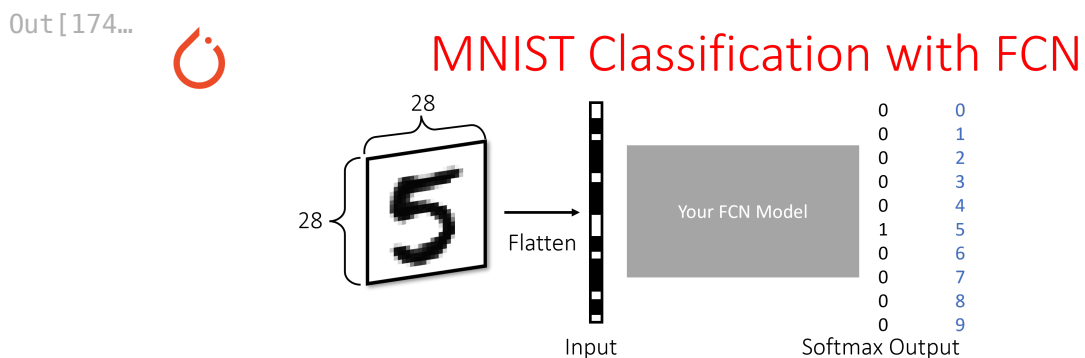
```
In [172... # Import necessary packages

%matplotlib inline

import matplotlib.pyplot as plt
import tqdm
import torch
import torchvision
import numpy as np
from sklearn.preprocessing import StandardScaler, MinMaxScaler

In [173... from IPython.display import Image # For displaying images in colab jupyter c

In [174... Image('lab3_exercise.png', width = 1000)
```



In this exercise, you will classify handwritten digits (28 x 28) using your own Fully Connected Network Architecture.

Prior to training your neural net, 1) Flatten each digit into 1D array of size 784, 2) Normalize the dataset using standard scaler and 3) Split the dataset into train/validation/test.

Design your own neural net architecture with your choice of hidden layers, activation functions, optimization method etc.

Your goal is to achieve a testing accuracy of >90%, with no restrictions on epochs.

Demonstrate the performance of your model via plotting the training loss, validation accuracy and printing out the testing accuracy.

Plot the testing samples where your model failed to classify correctly and print your model's best guess for each of them

44

## Prepare Data

```
In [175... # Load MNIST Dataset in Numpy

# 1000 training samples where each sample feature is a greyscale image with
# 1000 training targets where each target is an integer indicating the true
mnist_train_features = np.load('mnist_train_features.npy')
```

```
mnist_train_targets = np.load('mnist_train_targets.npy')

# 100 testing samples + targets
mnist_test_features = np.load('mnist_test_features.npy')
mnist_test_targets = np.load('mnist_test_targets.npy')

# Print the dimensions of training sample features/targets
print(mnist_train_features.shape, mnist_train_targets.shape)
# Print the dimensions of testing sample features/targets
print(mnist_test_features.shape, mnist_test_targets.shape)

print(mnist_train_targets[:10]) # print a sample
```

```
(1000, 28, 28) (1000,)
(100, 28, 28) (100,)
[5 0 4 1 9 2 1 3 1 4]
```

```
In [176... # Let's visualize some training samples

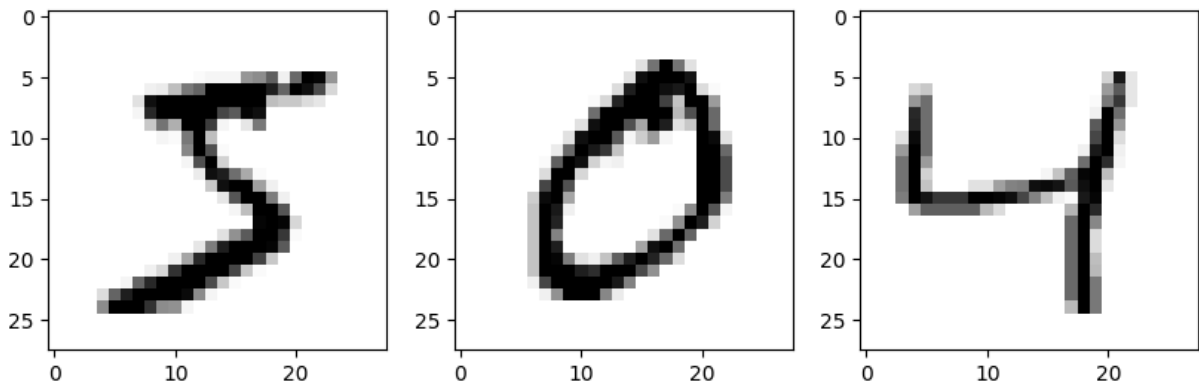
plt.figure(figsize = (10, 10))

plt.subplot(1,3,1)
plt.imshow(mnist_train_features[0], cmap = 'Greys')

plt.subplot(1,3,2)
plt.imshow(mnist_train_features[1], cmap = 'Greys')

plt.subplot(1,3,3)
plt.imshow(mnist_train_features[2], cmap = 'Greys')
```

```
Out[176... <matplotlib.image.AxesImage at 0x724f0f77a540>
```



```
In [177... # Reshape features via flattening the images
# This refers to reshape each sample from a 2d array to a 1d array.
# hint: np.reshape() function could be useful here

mnist_train_features = np.reshape(mnist_train_features, (1000, 784), copy=False)
mnist_test_features = np.reshape(mnist_test_features, (100, 784), copy=False)

print(mnist_train_features.shape, mnist_test_features.shape) # check the dim

(1000, 784) (100, 784)
```

```
In [178... # Scale the dataset according to standard scaling
def scaling(train, test, type = 'standard'): # scale using different methods
    if type == 'standard':
        scaler = StandardScaler()
        sc_train = scaler.fit_transform(train)
        sc_test = scaler.transform(test)
    elif type == 'minmax':
        scaler = MinMaxScaler()
        sc_train = scaler.fit_transform(train)
        sc_test = scaler.transform(test)
    else:
        raise ValueError("Invalid scaling type. Choose 'standard' or 'minmax'")
    return sc_train, sc_test

mnist_train_features, mnist_test_features = scaling(mnist_train_features, mnist_test_features,
                                                    type = 'standard')
```

```
In [179... # Split training dataset into Train (90%), Validation (10%)

valnum = int(0.1 * mnist_train_features.shape[0]) # 10% of training samples

mnist_validation_features = mnist_train_features[:valnum]
mnist_validation_targets = mnist_train_targets[:valnum]

mnist_train_features = mnist_train_features[valnum:]
mnist_train_targets = mnist_train_targets[valnum:]
```

## Define Model

```
In [180... class mnistClassification(torch.nn.Module):
    # added some arguments to make the model more flexible so that I could
    # quickly experiment with different architectures
    def __init__(self, input_dim, output_dim, layers=[128],
                  dropout_prob = 0.3):

        super(mnistClassification, self).__init__()
        self.num_layers = len(layers)
        # need different layer sizes for different number of layers so
        # that they are compatible
        if self.num_layers == 1:
            self.fc_in = torch.nn.Linear(input_dim, layers[0])
            self.fc_out = torch.nn.Linear(layers[0], output_dim)
            self.bn1 = torch.nn.BatchNorm1d(layers[0])
            self.drop1 = torch.nn.Dropout(dropout_prob)
        elif self.num_layers == 2:
            self.fc_in = torch.nn.Linear(input_dim, layers[0])
            self.fc1 = torch.nn.Linear(layers[0], layers[1])
            self.fc_out = torch.nn.Linear(layers[1], output_dim)
            self.bn1 = torch.nn.BatchNorm1d(layers[0])
            self.bn2 = torch.nn.BatchNorm1d(layers[1])
            self.drop1 = torch.nn.Dropout(dropout_prob)
            self.drop2 = torch.nn.Dropout(dropout_prob)
        elif self.num_layers == 3:
            self.fc_in = torch.nn.Linear(input_dim, layers[0])
```

```

        self.fc1 = torch.nn.Linear(layers[0], layers[1])
        self.fc2 = torch.nn.Linear(layers[1], layers[2])
        self.fc_out = torch.nn.Linear(layers[2], output_dim)
        self.bn1 = torch.nn.BatchNorm1d(layers[0])
        self.bn2 = torch.nn.BatchNorm1d(layers[1])
        self.bn3 = torch.nn.BatchNorm1d(layers[2])
        self.drop1 = torch.nn.Dropout(dropout_prob)
        self.drop2 = torch.nn.Dropout(dropout_prob)
        self.drop3 = torch.nn.Dropout(dropout_prob)
    else:
        raise ValueError('Invalid depth. Max depth is 3')
    # can choose activation function here
    self.activation = torch.nn.ReLU()

    # Foward pass defined as usual, put all in one line for smaller code
    # cell but the structure is as follows:
    # 1. input layer
    # 2. batch normalization
    # 3. dropout
    # 4. activation function
    # repeat for al layers
    # 5. output layer doesnt get batch normalization, dropout or
    # activation function applied

    # Note: We do not apply softmax to the output layer because
    # pytorch's CrossEntropyLoss function already applies softmax
    # internally and expects raw logits as input.
    # If we applied softmax here it would get applied twice,
    # resulting in instabilities and slower convergence.
    def forward(self, x):
        if self.num_layers == 1:
            out = self.fc_out(self.activation(self.drop1(self.bn1(
                self.fc_in(x)))))
        elif self.num_layers == 2:
            out = self.fc_out(self.activation(self.drop2(self.bn2(
                self.fc1(self.activation(self.drop1(self.bn1(
                    self.fc_in(x))))))))))
        elif self.num_layers == 3:
            out = self.fc_out(self.activation(self.drop3(self.bn3(
                self.fc2(self.activation(self.drop2(self.bn2(
                    self.fc1(self.activation(self.drop1(self.bn1(self.fc_in(
            else:
                raise ValueError('Invalid depth. Max depth is 3')
        return out

```

```

In [181]... # trying out initialization to improve performance
def initialize_weights_relu(m, activation='relu'):
    if isinstance(m, torch.nn.Linear): # Apply only to Linear layers
        # He initialization is best for ReLU
        if activation == 'relu':
            torch.nn.init.kaiming_normal_(m.weight, nonlinearity='relu')
        elif activation == 'tanh':
            torch.nn.init.kaiming_normal_(m.weight, nonlinearity='tanh')
        # Xavier initialization is best for sigmoid
        elif activation == 'sigmoid':
            torch.nn.init.xavier_normal_(m.weight, gain=torch.nn.init.calcul

```

```

if m.bias is not None:
    torch.nn.init.zeros_(m.bias) # Just biases initialized to zero

```

## Define Hyperparameters

```

In [182... # Initialize our neural network model with input and output dimensions
# using funnel architecture here
model = mnistClassification(input_dim=784, output_dim=10,
                             layers=[512, 128, 64], dropout_prob=0.3)

model.apply(lambda m: initialize_weights_relu(m, activation='relu')) # init

# Define the learning rate and epoch
learning_rate = 0.003
epochs = 150
print(f'Learning rate: {learning_rate}')
print(f'Epochs: {epochs}')

# Define the L2 regularization lambda parameter
reg_lambda = 0.06
print(f'Regularization lambda: {reg_lambda}')

# Define the mini batch size
batch_pwr = 5
batchsize = 2**batch_pwr
print(f'Batch size: {batchsize}')

# Define loss function and optimizer
# Optimizer is SGD with momentum 0.7, and learning rate and weight decay as
# defined above. Wanted to try without using Adam for developing better
# understanding of manual tuning
loss_func = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate,
                              weight_decay=reg_lambda, momentum=0.7)

# Run this line if you have PyTorch GPU version
if torch.cuda.is_available():
    model.cuda()

model

```

```

Learning rate: 0.003
Epochs: 150
Regularization lambda: 0.06
Batch size: 32

```

```

Out[182... mnistClassification(
    (fc_in): Linear(in_features=784, out_features=512, bias=True)
    (fc1): Linear(in_features=512, out_features=128, bias=True)
    (fc2): Linear(in_features=128, out_features=64, bias=True)
    (fc_out): Linear(in_features=64, out_features=10, bias=True)
    (bn1): BatchNorm1d(512, eps=1e-05, momentum=0.1, affine=True, track_runni
ng_stats=True)
    (bn2): BatchNorm1d(128, eps=1e-05, momentum=0.1, affine=True, track_runni
ng_stats=True)
    (bn3): BatchNorm1d(64, eps=1e-05, momentum=0.1, affine=True, track_runnin
g_stats=True)
    (drop1): Dropout(p=0.3, inplace=False)
    (drop2): Dropout(p=0.3, inplace=False)
    (drop3): Dropout(p=0.3, inplace=False)
    (activation): ReLU()
)

```

## Identify Tracked Values

```

In [183... # Placeholders for training loss and validation accuracy during training
# Training loss should be tracked for each iteration
# (1 iteration -> single forward pass to the network)
# Validation accuracy should be evaluated every 'Epoch'
# (1 epoch -> full training dataset)
# If using batch gradient, 1 iteration = 1 epoch

train_loss_list = []
validation_accuracy_list = []
# also want to look at validation loss to check for overfitting
validation_loss_list = []

```

## Train Model

```

In [184... # Convert the training, validation, testing dataset (NumPy arrays) into torch

train_inputs = torch.from_numpy(mnist_train_features).float()
train_targets = torch.from_numpy(mnist_train_targets).long()
#idk why we have to use long here but it works

validation_inputs = torch.from_numpy(mnist_validation_features).float()
validation_targets = torch.from_numpy(mnist_validation_targets).long()

testing_inputs = torch.from_numpy(mnist_test_features).float()
testing_targets = torch.from_numpy(mnist_test_targets).long()

# Training Loop -----

for epoch in tqdm.trange(epochs):
    # get random indices to shuffle the training data
    indices = torch.randperm(train_inputs.size(0))

    # Shuffle both tensors while preserving the correspondence between
    # inputs and targets

```

```

temp_train_inputs = train_inputs[indices]
temp_train_targets = train_targets[indices]

# Iterate over mini-batches of training data
for i in range(0, len(train_inputs), batchsize):
    # indexing doesnt error here cause exclusivity, we dont waste any
    # remainder of data here cause of the way python handles slicing
    mb_inputs = temp_train_inputs[i:i+batchsize]
    mb_targets = temp_train_targets[i:i+batchsize]

    optimizer.zero_grad()

    train_outputs = model(mb_inputs)

    loss = loss_func(train_outputs, mb_targets)

    train_loss_list.append(loss.item())

    loss.backward()

    optimizer.step()

# Compute Validation Accuracy -----
model.eval() # turn off dropout and batch normalization
with torch.no_grad(): # turn off gradient tracking

    validation_outputs = model(validation_inputs)

    val_loss = loss_func(validation_outputs, validation_targets)
    # append validation loss for each epoch, not minibatch
    validation_loss_list.append(val_loss.item())

    predicted = torch.argmax(validation_outputs, dim=1)
    # dont even need softmax here cause we aren't feeding it to
    # cross entropy. Raw logit indices correspond to labels 1-10

    # Compute the accuracy
    correct = (predicted == validation_targets).sum().item()
    # boolean logic returns true false array
    # since True=1 False=0 we can just sum
    accuracy = correct / validation_targets.size(0)
    # divide by number of samples to get accuracy

    validation_accuracy_list.append(accuracy) # save accuracy for plots
model.train() # turn on dropout and batch normalization for next iteration

```

100%|██████████| 150/150 [00:07<00:00, 20.25it/s]

```

In [185... # Figuring out array sizes so that I can plot the loss per minibatch and
# the validation loss per epoch superimposed on same graph
datasize = len(train_inputs)
print(f'Training data size: {datasize} \nBatchsize: {batchsize} \nEpochs: {e
#print(len(train_inputs)/batchsize)

# +1 accounts for remainder slice iteration in each minibatch

```

```

data_divs = (datasize // batchsize) + 1
print(f'Datasize/batchsize, plus remainder slice --> # of training data divi
step_range = np.arange(0, data_divs*epochs)
# define "xaxis" for minibatch training loss
epoch_spacing_range = np.linspace(0, epochs*data_divs, epochs+1)[1:]
# define "xaxis" for validation accuracy

#print(step_range)
print(epoch_spacing_range[-1])
# Epoch loss should be plotted at the end of each epoch, not end of each
# minibatch. Should end at 5699 but this is good enough for diagnostic
# purposes

print(f'Training loss list length:          {len(train_loss_list)}')
print(f'step_range length:                  {len(step_range)}')

print(f'Validation accuracy list length:    {len(validation_accuracy_list)}')
print(f'Validation loss list length:        {len(validation_loss_list)}')
print(f'epoch_spacing_range length:        {len(epoch_spacing_range)}')

```

Training data size: 900

Batchsize: 32

Epochs: 150

Datasize/batchsize, plus remainder slice --> # of training data divisions by batch in each epoch: 29

4350.0

Training loss list length: 4350

step\_range length: 4350

Validation accuracy list length: 150

Validation loss list length: 150

epoch\_spacing\_range length: 150

## Visualize and Evaluate Model

In [186... *# Import seaborn for prettier plots*

```
import seaborn as sns
```

In [187... *# Visualize training loss*

```

plt.figure(figsize = (12, 6))

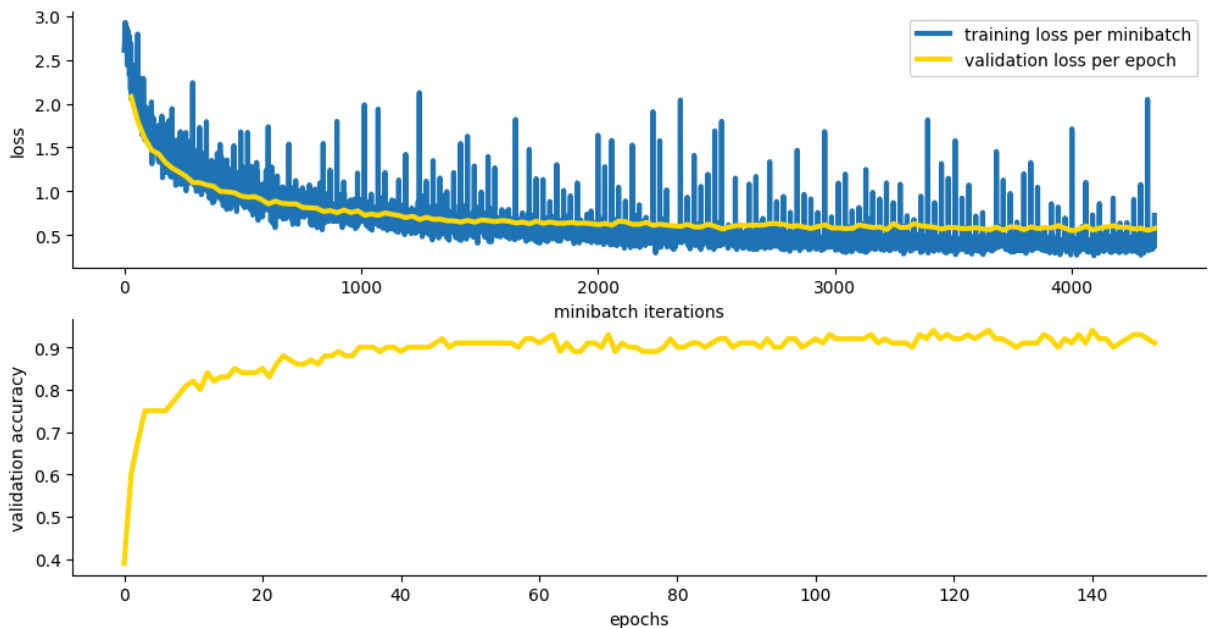
# Visualize training loss with respect to iterations
# (1 iteration -> single mini batch)
plt.subplot(2, 1, 1)
# this allows for visualizing minibatch training loss noise but noiseless
# validation loss
# use x axes computed in prev cell
plt.plot(step_range, train_loss_list, linewidth = 3)
plt.plot(epoch_spacing_range, validation_loss_list, linewidth = 3, color = 'r')
plt.ylabel("loss")
plt.xlabel("minibatch iterations")
plt.legend(['training loss per minibatch', 'validation loss per epoch'])

```



```
sns.despine()

# Visualize validation accuracy with respect to epochs
plt.subplot(2, 1, 2)
plt.plot(validation_accuracy_list, linewidth = 3, color = 'gold')
plt.ylabel("validation accuracy")
plt.xlabel("epochs")
sns.despine()
```



```
In [188... # Compute the testing accuracy
model.eval()
with torch.no_grad():

    test_outputs = model(testing_inputs) # raw logits
    test_predicted = torch.argmax(test_outputs, dim=1)
    # convert to predicted labels as before

    # Compute the correct/ incorrect predictions and keep in array
    # with corresponding indices
    test_sucesses = test_predicted == testing_targets

    test_correct_num = (test_sucesses).sum().item()
    test_accuracy = test_correct_num / testing_targets.size(0)

    print(f"Test Accuracy: {test_accuracy*100}%")
model.train()
```

Test Accuracy: 94.0%

```

Out[188... mnistClassification(
    (fc_in): Linear(in_features=784, out_features=512, bias=True)
    (fc1): Linear(in_features=512, out_features=128, bias=True)
    (fc2): Linear(in_features=128, out_features=64, bias=True)
    (fc_out): Linear(in_features=64, out_features=10, bias=True)
    (bn1): BatchNorm1d(512, eps=1e-05, momentum=0.1, affine=True, track_runni
ng_stats=True)
    (bn2): BatchNorm1d(128, eps=1e-05, momentum=0.1, affine=True, track_runni
ng_stats=True)
    (bn3): BatchNorm1d(64, eps=1e-05, momentum=0.1, affine=True, track_runnin
g_stats=True)
    (drop1): Dropout(p=0.3, inplace=False)
    (drop2): Dropout(p=0.3, inplace=False)
    (drop3): Dropout(p=0.3, inplace=False)
    (activation): ReLU()
)

```

```

In [189... # Plot 5 incorrectly classified testing samples and print the model predicti
# You can use np.reshape() to convert flattened 1D array back to 2D array

failed_i = np.where(test_sucesses == False)[0]
# use boolean array from prev cell to find failed indices
print(failed_i[:10])

```

```

In [190... # convert back to 2D
mnist_test_features_resized = np.reshape(mnist_test_features,
                                           (100, 28, 28), copy=False)

plt.figure(figsize = (10, 10))

plt.subplot(3,5,1)
# show the first failed sample
plt.imshow(mnist_test_features_resized[failed_i[0]], cmap = 'Greys')
# show the predicted and true labels for the failed sample
plt.title("Predicted: " + str(test_predicted[failed_i[0]].item()) +
          "\nTrue: " + str(testing_targets[failed_i[0]].item()))

plt.subplot(3,5,2)
plt.imshow(mnist_test_features_resized[failed_i[1]], cmap = 'Greys')
plt.title("Predicted: " + str(test_predicted[failed_i[1]].item()) +
          "\nTrue: " + str(testing_targets[failed_i[1]].item()))

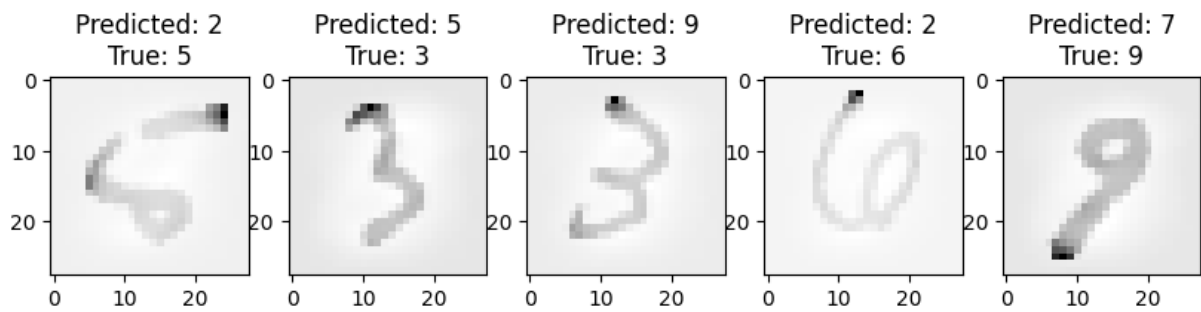
plt.subplot(3,5,3)
plt.imshow(mnist_test_features_resized[failed_i[2]], cmap = 'Greys')
plt.title("Predicted: " + str(test_predicted[failed_i[2]].item()) + "\nTrue:

plt.subplot(3,5,4)
plt.imshow(mnist_test_features_resized[failed_i[3]], cmap = 'Greys')
plt.title("Predicted: " + str(test_predicted[failed_i[3]].item()) + "\nTrue:

plt.subplot(3,5,5)
plt.imshow(mnist_test_features_resized[failed_i[4]], cmap = 'Greys')
plt.title("Predicted: " + str(test_predicted[failed_i[4]].item()) + "\nTrue:

```

```
Out[190... Text(0.5, 1.0, 'Predicted: 7\nTrue: 9')
```



```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```